

# AI security review that never sees the cloud.

Cloud AI code review means uploading your proprietary source to someone else's servers. For regulated, air-gapped, or simply confidential codebases, that's a **non-starter**.


Lemonade-Sec runs the **entire** pipeline — static analysis *and* AI reasoning — **locally, on your own machine**, powered by AMD Lemonade.

**0**BYTES LEAVE THE  
BOX**35**RULES · WEB2 +  
WEB3**\$0**

MIT CORE · NO KEYS

## ARCHITECTURAL FACT · WHERE YOUR CODE GOES

Private by design: source never leaves the machine



# Scan → local-AI triage → report, in one command.

## The pipeline

- **SAST core** — 35 regex rules across secrets, injection, SSRF, XSS, crypto, JWT + **web3/Solidity** (reentrancy, arbitrary delegatecall...).
- **Local AI triage** — each finding + code context + attack playbook → your **Lemonade** model returns verdict, confidence, and a PoC idea.
- **OpenAI-compatible** — auto-discovers the model via `/models`; works with any local server.
- **Graceful offline fallback** — no server? A built-in heuristic still produces a report. **Never a network call you didn't ask for.**

**0**

PIP DEPENDENCIES

**14**FINDINGS ON  
DEMO**1 cmd**


TO RUN

Stdlib-only · authorized-use gate (`--scope`) · findings are candidates

Coverage across web2 + web3 in one pass

MEASURED · `python3 lemonade_sec.py samples --triage`

14 findings on the demo sample, each AI-triaged locally



# Two waves collide: local AI + AI-native AppSec.

## Who pays — and can't use cloud AI

- Banks & fintech (source-code residency rules)
- Defense, gov, critical infrastructure (air-gapped)
- Healthcare & anyone under strict IP/NDA
- Enterprises banning code in public chatbots

## Why now

- On-device LLMs went usable: **Lemonade 2.7k★**, AMD Strix Halo / Ryzen AI, 128 GB unified memory.
- AppSec is shifting to AI triage — but privacy blocks cloud adoption.

## Market sizing (modeled, directional)

TAM · APPSEC / SAST + AI CODE TOOLS

≈ **\$6–8B**

2026, growing double-digits

SAM · PRIVACY-CONSTRAINED / ON-PREM-REQUIRED ORGS

≈ **\$0.9B**

regulated + air-gapped segment

SOM · 3-YR OBTAINABLE (OPEN-CORE WEDGE)

≈ **\$15–25M**

land free → expand to paid teams


Figures are modeled from public market estimates — directional sizing, not audited.

# Free core drives adoption; teams pay for the enterprise layer.

| TIER        | PRICE (MODELED) | WHAT'S IN IT                                                          |
|-------------|-----------------|-----------------------------------------------------------------------|
| Core<br>MIT | \$0             | CLI, 35 rules, local Lemonade triage, offline fallback                |
| Team        | \$15 /dev/mo    | SARIF + CI gates, policy-as-code, baseline/diff, dashboards           |
| Enterprise  | \$40 /dev/mo    | SSO/RBAC, audit, air-gap deploy, managed local-model ops, support/SLA |

## Unit economics (illustrative)

- Land free (OSS) → expand to Team → Enterprise.  
NRR>120% target via seat growth.
- Near-zero COGS: inference runs on the **customer's** hardware, not ours.
- 200-dev enterprise ≈ **\$96k ARR**; 30 such logos ≈ **\$2.9M ARR**.



Incumbents charge per seat and require the cloud. Our core is free and local; we monetize CI/governance/support — not the scan.

# Built, tested, and riding the AMD local-AI wave.

## Traction today **REAL**

- Working tool: scans + local-AI triage, **14 findings** on the demo, MIT on GitHub.
- **Measured on AMD Instinct MI300X** (torch-ROCM, fp16): **64 tok/s** decode · 13.8k prefill · **19 ms** TTFT — the triage pipeline runs on real AMD silicon.
- Submitted to the **AMD Lemonade Developer Challenge**; author is an **AMD ACT I winner** · **Developer Featured** — warm ecosystem channel.

## Roadmap

- SARIF + GitHub code-scanning · AST taint-tracking to cut false positives
- Ship as a Lemonade Marketplace app · design-partner pilots

## Valuation frame (illustrative · pre-revenue)

Comparables in AI-native AppSec: **Semgrep**, **Snyk**, **Socket** — venture-backed on OSS-led adoption.

MODELED EARLY-STAGE RANGE

**\$3–6M**

pre-seed, on team + wedge + AMD tailwind — a frame, not a quote

## The ask

Design partners in regulated eng orgs · a Lemonade marketplace slot · and the AMD hardware/relationship to make local-AI AppSec the default.